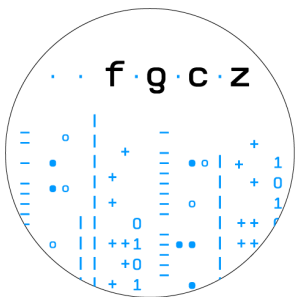


Mapping Reads

Dr. Hubert Rehrauer





Purpose of Read Mapping

- Identify the origin of a read
- RNA-seq: Infer the gene that was expressed and that generated the read
- Is it always possible to find the read origin?

Sequence Alignment

- Optimization problem: Find the alignment with the highest score
- Global alignment (end-to-end): Needleman-Wunsch
- Local alignment: Smith-Waterman

Score matrix:

	a	c	g	t
a	+2	-2	-1	-2
c	-2	+2	-2	-1
g	-1	-2	+2	-2
t	-2	-1	-2	+2

Gap score: -3

c	c	t	a	g	t	t	c	a	t
				x		x			
t	g	t	a	a	t	-	c	a	c
		+2	+2	-1	+2	-3	+2	+2	

Alignment score = 6



Local alignments (Smith-Waterman)

Generally useful for local alignments, determining regions of similarity between two strings (here, DNA sequence)

Dynamic programming based algorithm: gives optimal alignment, with respect to scoring system

Need to set scores for **match**, penalties for **mismatch** and **gap** (typically, mismatch penalties are set according to evolutionary knowledge)

Match: +1; Mismatch: -1; Gap: -2

	.	A	A	T	G	T
.	0	0	0	0	0	0
A	0					
T	0					
G	0					
A	0					
C	0					

$$H(i, j) = \max \left\{ \begin{array}{l} 0 \\ H(i-1, j-1) + w(a_i, b_j) \\ H(i-1, j) + w(a_i, -) \\ H(i, j-1) + w(-, b_j) \end{array} \right\}, \quad 1 \leq i \leq m, 1 \leq j \leq n$$

Match/Mismatch

Deletion

Insertion



Local alignments (Smith-Waterman)

Generally useful for local alignments, determining regions of similarity between two strings (here, DNA sequence)

Dynamic programming based algorithm: gives optimal alignment, with respect to scoring system

Need to set scores for **match**, penalties for **mismatch** and **gap**

Match: +1; Mismatch: -1; Gap: -2

	.	A	A	T	G	T
.	0	0	0	0	0	0
A	0	1	1			
T	0					
G	0					
A	0					
C	0					

$$H(i, j) = \max \left\{ \begin{array}{l} 0 \\ H(i-1, j-1) + w(a_i, b_j) \\ H(i-1, j) + w(a_i, -) \\ H(i, j-1) + w(-, b_j) \end{array} \right\}, \quad 1 \leq i \leq m, 1 \leq j \leq n$$

Match/Mismatch

Deletion

Insertion



Local alignments (Smith-Waterman)

Generally useful for local alignments, determining regions of similarity between two strings (here, DNA sequence)

Dynamic programming based algorithm: gives optimal alignment, with respect to scoring system

Need to set scores for **match**, penalties for **mismatch** and **gap**

Match: +1; Mismatch: -1; Gap: -2

	.	A	A	T	G	T
.	0	0	0	0	0	0
A	0	1	1	0	0	0
T	0	0	0	2	0	1
G	0					
A	0					
C	0					

$$H(i, j) = \max \left\{ \begin{array}{l} 0 \\ H(i-1, j-1) + w(a_i, b_j) \\ H(i-1, j) + w(a_i, -) \\ H(i, j-1) + w(-, b_j) \end{array} \right\}, \quad 1 \leq i \leq m, 1 \leq j \leq n$$

Match/Mismatch

Deletion

Insertion



Local alignments (Smith-Waterman)

Generally useful for local alignments, determining regions of similarity between two strings (here, DNA sequence)

Dynamic programming based algorithm: gives optimal alignment, with respect to scoring system

Need to set scores for **match**, penalties for **mismatch** and **gap**

Match: +1; Mismatch: -1; Gap: -2

	.	A	A	T	G	T
.	0	0	0	0	0	0
A	0	1	1	0	0	0
T	0	0	0	2	0	1
G	0	0	0	0	3	1
A	0	1	1	0	1	2
C	0	0	0	0	0	0

$$H(i, j) = \max \left\{ \begin{array}{l} 0 \\ H(i-1, j-1) + w(a_i, b_j) \\ H(i-1, j) + w(a_i, -) \\ H(i, j-1) + w(-, b_j) \end{array} \right\}, \quad 1 \leq i \leq m, 1 \leq j \leq n$$

Match/Mismatch

Deletion

Insertion



Local alignments (Smith-Waterman)

Traceback:

Start with maximum score (here, 3)

Follow path that gives multiple score

	.	A	A	T	G	T
.	0	0	0	0	0	0
A	0	1	1	0	0	0
T	0	0	0	2	0	1
G	0	0	0	0	3	1
A	0	1	1	0	1	2
C	0	0	0	0	0	0

$$H(i, j) = \max \begin{cases} 0 \\ H(i-1, j-1) + w(a_i, b_j) \\ H(i-1, j) + w(a_i, -) \\ H(i, j-1) + w(-, b_j) \end{cases}, 1 \leq i \leq m, 1 \leq j \leq n$$

Match/Mismatch

Deletion

Insertion



Local alignments (Smith-Waterman)

Traceback:


Start with maximum score (here, 3)

Follow path that gives multiple score

This gives:

AATGT-

-ATGAC



$$H(i, j) = \max \begin{cases} 0 \\ H(i-1, j-1) + w(a_i, b_j) & \text{Match/Mismatch} \\ H(i-1, j) + w(a_i, -) & \text{Deletion} \\ H(i, j-1) + w(-, b_j) & \text{Insertion} \end{cases}, \quad 1 \leq i \leq m, 1 \leq j \leq n$$



Local alignments (Smith-Waterman)

Other considerations:

Natural extension is to have a different *gap opening* and a *gap extension* penalty (former generally being larger)

GO penalty=-2; GE penalty=-1; Mismatch=-1; Match=1

With above penalties, which is the best scoring alignment?

AT-C-GT

ATC--GT

AT-C--GT

ATTTTGT

ATTTTGT

ATT-TTGT




Alignment Tools

- Dynamic programming is slow
- Speed-up with heuristics
 - e.g. exactly align short subsequences and extend these alignments
 - e.g. BLAST / BLAT
 - no longer guaranteed to find the best alignments
 - exact matches are found by index lookup


Index Genome: Spaced Seeds

- Tags and tag-sized pieces of reference are cut into small “seeds.”
- Pairs of spaced seeds are stored in an index.
- Look up spaced seeds for each tag.
- For each “hit,” confirm the remaining positions.



Index Genome: Burrows-Wheeler Transform

- Store entire reference genome.
- Align tag base by base from the end.
- When tag is traversed, all active locations are reported.
- If no match is found, then back up and try a substitution.




The Burrows-Wheeler transform (1994; 1983)

c a c a a c g \$



c	a	c	a	a	c	g	\$
a	c	a	a	c	g	\$	c
c	a	a	c	g	\$	c	a
a	a	c	g	\$	c	a	c
a	c	g	\$	c	a	c	a
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c
\$	c	a	c	a	a	c	g



\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c



g c c a a \$ a c

The “Last-First mapping” property

The relative ordering of a particular character (say **c**) in column 1 is the same as that in the last column

c_1 a c_2 a a c_3 g \$

\$	c_1	a	c_2	a	a	c_3	g
a	a	c_3	g	\$	c_1	a	c_2
a	c_2	a	a	c_3	g	\$	c_1
a	c_3	g	\$	c_1	a	c_2	a
c_2	a	a	c_3	g	\$	c_1	a
c_1	a	c_2	a	a	c_3	g	\$
c_3	g	\$	c_1	a	c_2	a	a
g	\$	c_1	a	c_2	a	a	c_3

The “Last-First mapping” property

The relative ordering of a particular character (say **c**) in column 1 is the same as that in the last column

c_1 a c_2 a a c_3 g \$

\$	c_1	a	c_2	a	a	c_3	g
a	a	c_3	g	\$	c_1	a	c_2
a	c_2	a	a	c_3	g	\$	c_1
a	c_3	g	\$	c_1	a	c_2	a
c_2	a	a	c_3	g	\$	c_1	a
c_1	a	c_2	a	a	c_3	g	\$
c_3	g	\$	c_1	a	c_2	a	a
g	\$	c_1	a	c_2	a	a	c_3

Proof:

Suppose cX and cY are cyclic permutations of the input T .

Suppose $cX < cY$ (in lexicographical ordering)

Then $Xc < Yc$ (in lexicographical ordering)

The LF-mapping property follows.

BWT is reversible

\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c

g c c a a \$ a c



\$ a a a c c c g

BWT is reversible

\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c

g\$ ca ca aa ac \$c ac cg



\$c aa ac ac ca ca cg g\$

BWT is reversible

\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c

Lookup AAC

\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c

Range \leftarrow range of last character in 1st column

While characters left (and nonzero range):

Lookup first and last match to preceding character in final column

Range \leftarrow LF-mapping of first and last match

Lookup AAC

\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c

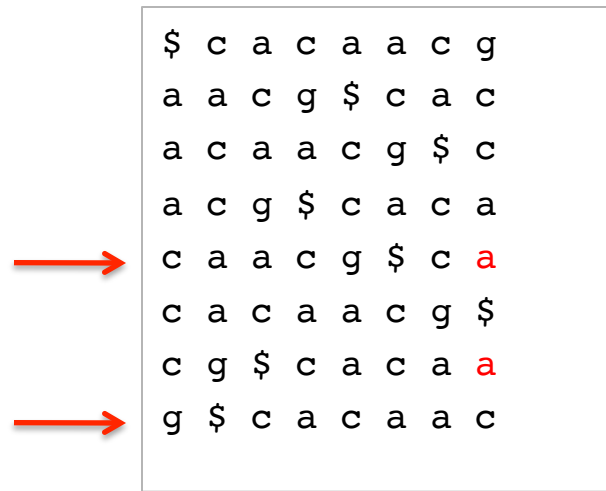
Range \leftarrow range of last character in 1st column

While characters left (and nonzero range):

Lookup first and last match to preceding character in final column

Range \leftarrow LF-mapping of first and last match

Lookup **A**C



\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c


Range \leftarrow range of last character in 1st column

While characters left (and nonzero range):

Lookup first and last match to preceding character in final column

Range \leftarrow LF-mapping of first and last match

Lookup **A**C



Range \leftarrow range of last character in 1st column

While characters left (and nonzero range):

Lookup first and last match to preceding character in final column

Range \leftarrow LF-mapping of first and last match

Lookup AAC

\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c

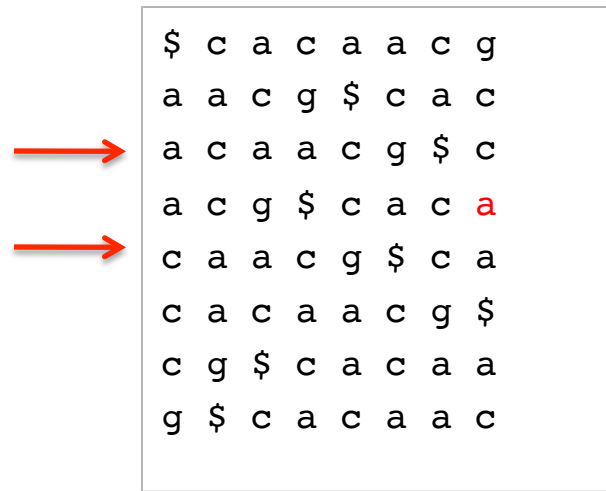
Range \leftarrow range of last character in 1st column

While characters left (and nonzero range):

Lookup first and last match to preceding character in final column

Range \leftarrow LF-mapping of first and last match

Lookup **A**AC



\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c


Range \leftarrow range of last character in 1st column

While characters left (and nonzero range):

Lookup first and last match to preceding character in final column

Range \leftarrow LF-mapping of first and last match

Lookup **A**AC



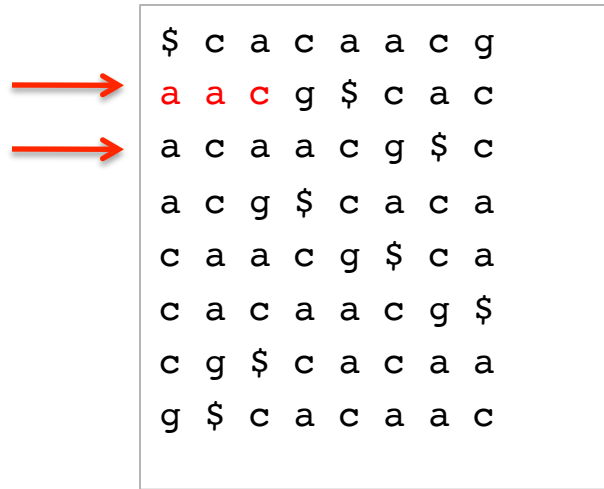
Range \leftarrow range of last character in 1st column

While characters left (and nonzero range):

Lookup first and last match to preceding character in final column

Range \leftarrow LF-mapping of first and last match

Lookup **A**AC



\$	c	a	c	a	a	c	g
a	a	c	g	\$	c	a	c
a	c	a	a	c	g	\$	c
a	c	g	\$	c	a	c	a
c	a	a	c	g	\$	c	a
c	a	c	a	a	c	g	\$
c	g	\$	c	a	c	a	a
g	\$	c	a	c	a	a	c

Range \leftarrow range of last character in 1st column

While characters left (and nonzero range):

Lookup first and last match to preceding character in final column

Range \leftarrow LF-mapping of first and last match



Comparison

Spaced seeds

- Requires ~50Gb of memory.
- Runs 30-fold slower.
- More straightforward to program.
- Examples:
 - MAQ
 - Shrimp
- More tolerant to
 - sequence variations
 - sequencing errors

Burrows-Wheeler

- Requires <2Gb of memory.
- Runs 30-fold faster.
- More complicated to program.
- Examples:
 - bowtie
 - BWA
 - tophat (uses bowtie)
 - STAR



Alignment with Mismatches

- Mismatches can occur because of
 - sequencing error (error rate $\sim 1/500$)
 - mutation; (human mutation rate $\sim 1/1e4$)
 - \rightarrow if reads are long ($> 100\text{nt}$) reads with mismatches will not be rare; more than $\sim 10\%$ of the reads may have a mismatch
- If there is a sequence mismatch the index lookup fails! How to find nevertheless an alignment?



BWT Alignment with Mismatches

- Strategies:
 - If the BWT lookup fails at position k , try all different bases at position k
 - drawback: computing effort grows exponentially with the number of mismatches
 - implemented e.g. in bowtie
 - Chop reads in segments (seeds) and align those mismatch-free and stitch seed alignments together
 - implemented e.g. in bowtie



Consideration: Sequencing Errors – PHRED quality

- Each called base is given a quality score Q
- $Q = 10 * \log_{10}(\text{estimated probability that call is wrong})$

10 prob = 0.1

20 prob = 0.01

30 prob = 0.001

[Q30 often used as a threshold for useful sequence data]

→ down-weight low-quality bases when computing the alignment score


- These quality scores are called PHRED scores. They were first introduced for capillary sequencers.
- PHRED scores are determined by the sequencer that directly rates how reliable the measured signal is



Considerations: Read trimming

- Sequencers have systematic errors
- Illumina sequencers have a higher error rate at the first few bases (1-5)
- Basically all sequencers have increasing error rate towards the end of the read
- Hard trimming: trim a fixed number of bases from the beginning and/or end
- Quality trimming: cut the end of the read as soon as the base quality drops below a threshold

Considerations: Read trimming



- If the inserted DNA/RNA fragment is too short, the read will contain part of the adapter
- Adapter trimming can be challenging if
 - if the insert is 90 – 100 bases
 - if the read has many sequencing errors



Multiple Alignments

- A read may have multiple valid alignments with identical or similarly good alignment scores
- Aligners allow to choose different reporting strategies:
 - Randomly select one alignment from the top-scoring alignments
 - Report all alignments that are within *delta* of the top-scoring alignment; clip if more than *Nmax* alignments are found
 - Report only alignments if they are unique (no other alignment within *delta* of the alignment score)
 - Do not report anything if more than *Nmax* valid alignments are found
 - ...
- Whether a read has a **unique alignment** depends on
 - the read sequence and the sequence homology of the organism
 - the search algorithm of the aligner
 - the reporting strategy of the aligner



Read Alignment Summary (I)

How to map billions of reads?

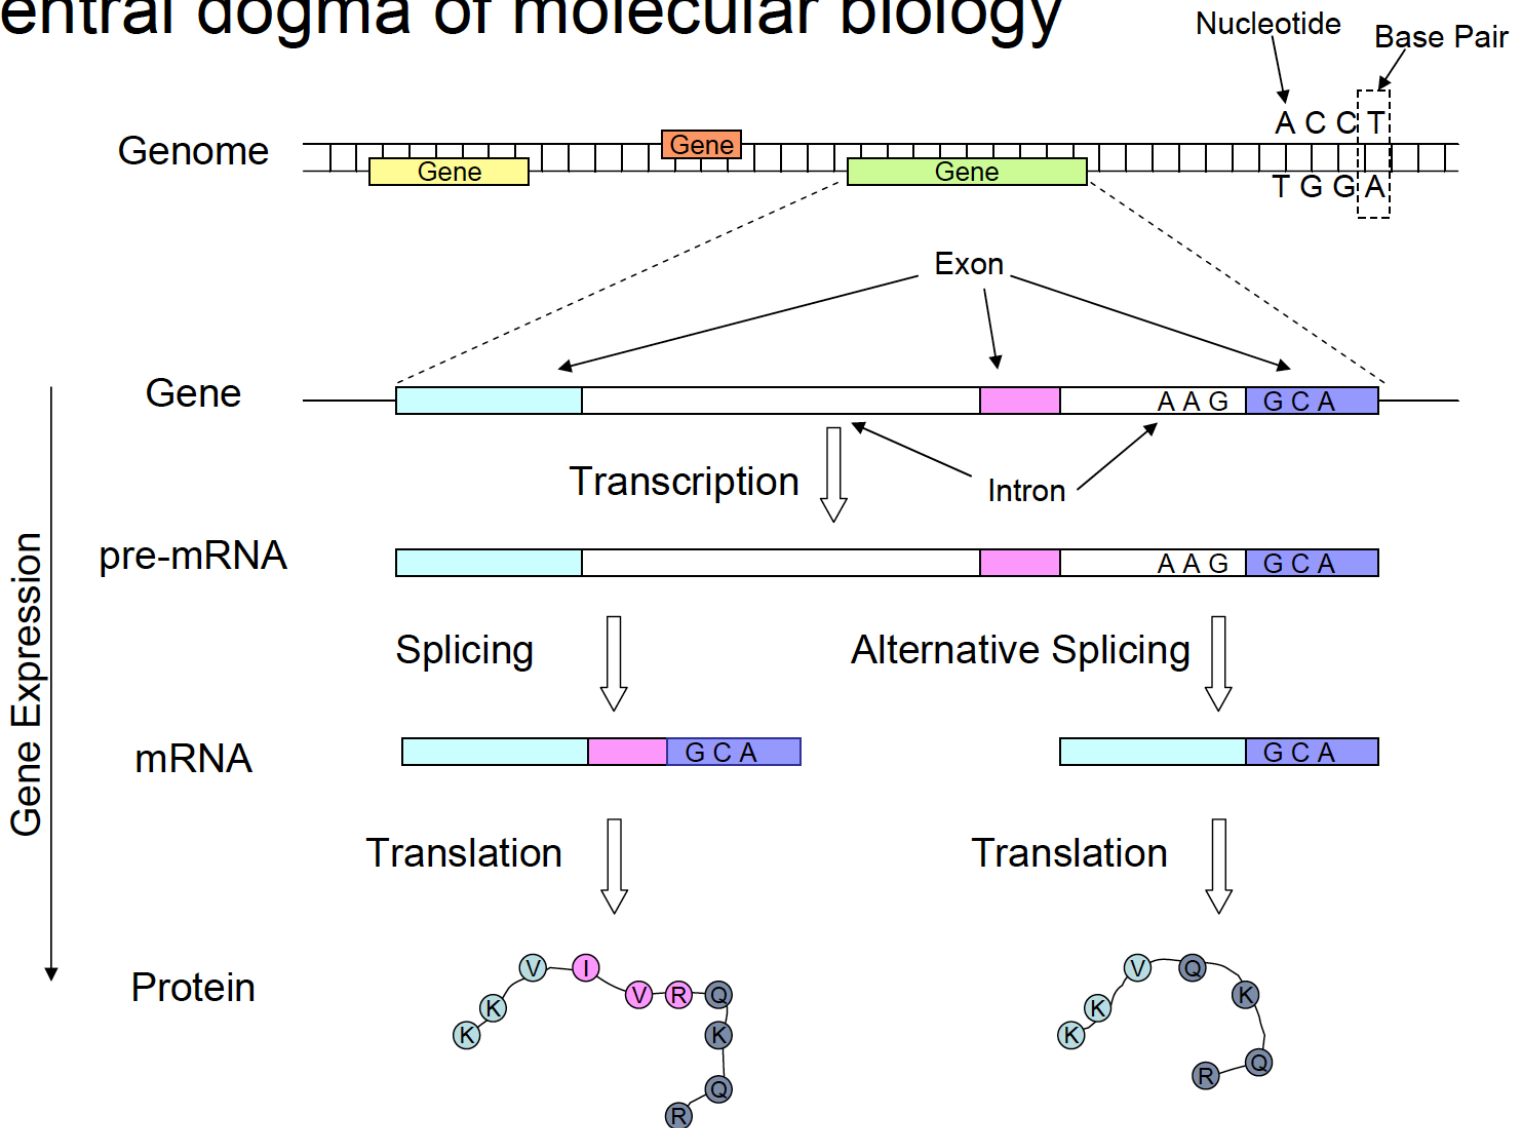
- The major aligners use the Burrows-Wheeler-Transform to create an index of the reference. Even for the human genome the index fits into 3GB RAM.
- Reads are aligned by **index lookup** not by sequence comparison
- The lookup of a perfect match read is faster than loading the read and writing the alignment coordinates to the output file!
- If the lookup fails because of SNPs or sequencing errors, an actual sequence comparison is done –and this can take time!



Read Alignment Summary (II)

- Alignment is a score optimization problem
- Aligners find the alignment with the **highest alignment score**
- Aligners do not run an extensive search, they use **index lookup and heuristics** to find the alignment with highest score
- If the heuristics are well chosen and the scores are well defined, then the alignment will be very often the correct alignment

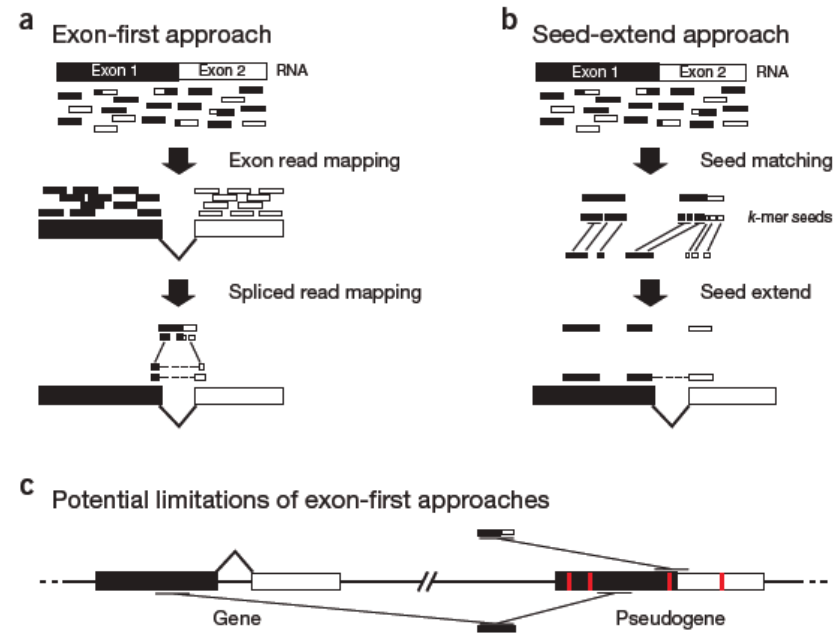
Central dogma of molecular biology



92–94% of human genes undergo alternative splicing,

RNA-seq Mapping

- Mapping targets:
 - transcriptome
 - splice junction library
 - genome
- Mouse retina 60+60bp reads:
 - 41 of 91 Mio map to junctions



Strategies:

- Exon-first: fast
- Seed-extend:
 - good with polymorphisms
 - simultaneous spliced/ unspliced mapping

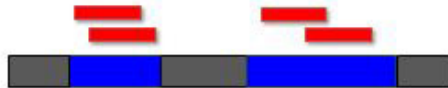
Tophat2 pipeline

(1) Transcriptome alignment (optional)



(2) Genome alignment

Reads spanning a single exon are mapped



Multi-exon spanning reads are unmapped



(3) Spliced alignment

Reads

Reads are aligned against transcriptome.

Transcriptome index

Reads are aligned against genome.

Genome index

(3) Spliced alignment

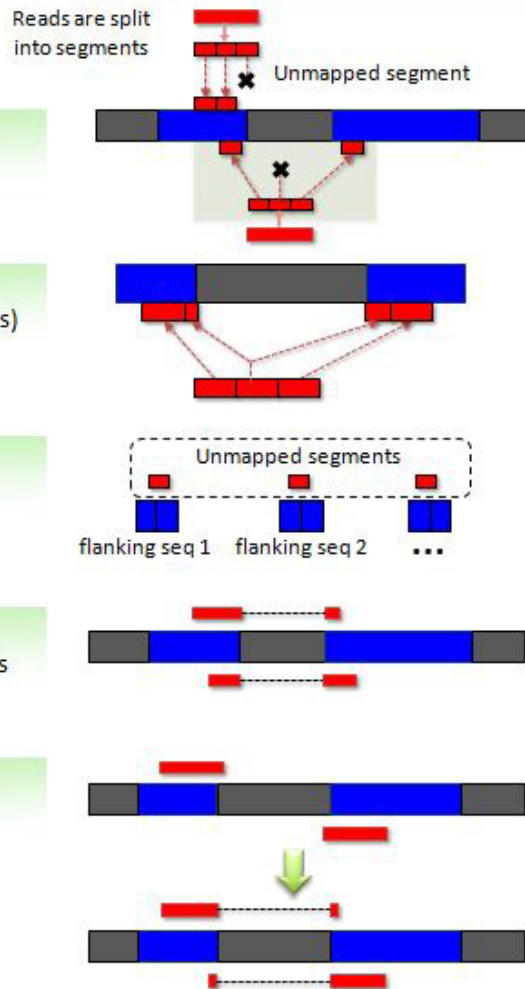
(3-1) Segment alignment to genome

(3-2) Identification of splice sites
(including indels and fusion break points)

(3-3) Segments aligned to junction
flanking sequences

(3-4) Segment alignments stitched
together to form whole read alignments

(3-5) Re-alignment of reads minimally
overlapping introns



Reads are split into smaller segments
which are then aligned to the genome.

Genome index

Segment mappings are used to find potential splice sites
usually when the distance between the mapped positions
of the left and the right segments are longer than the
length of the middle part of a read.

Sequences flanking a splice site are concatenated
and segments are aligned to them.


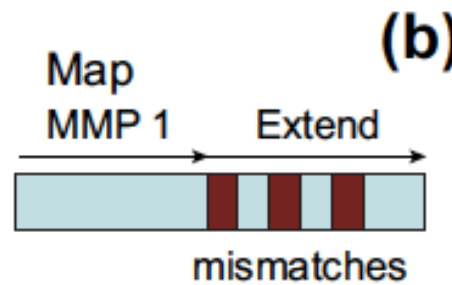
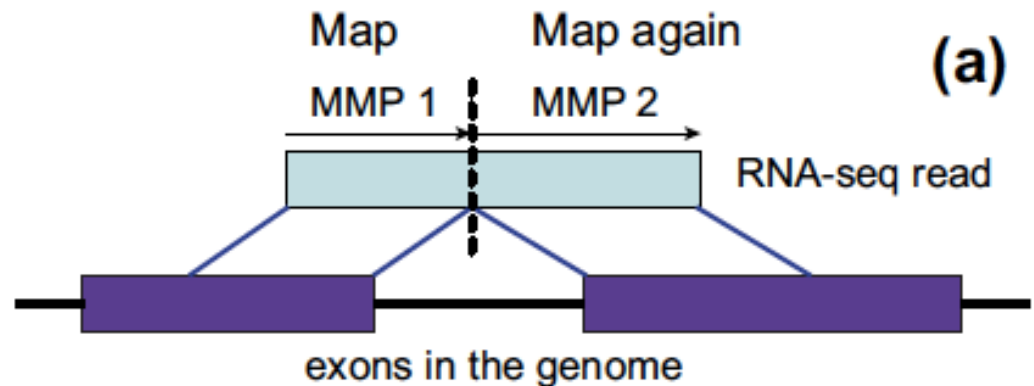
Junction flanking index

Mapped segments against either genome or flanking
sequences are gathered to produce whole read alignments.

Genome mapped reads with alignments extending a few
bases into introns are re-aligned to exons instead.

STAR: universal RNA-seq aligner

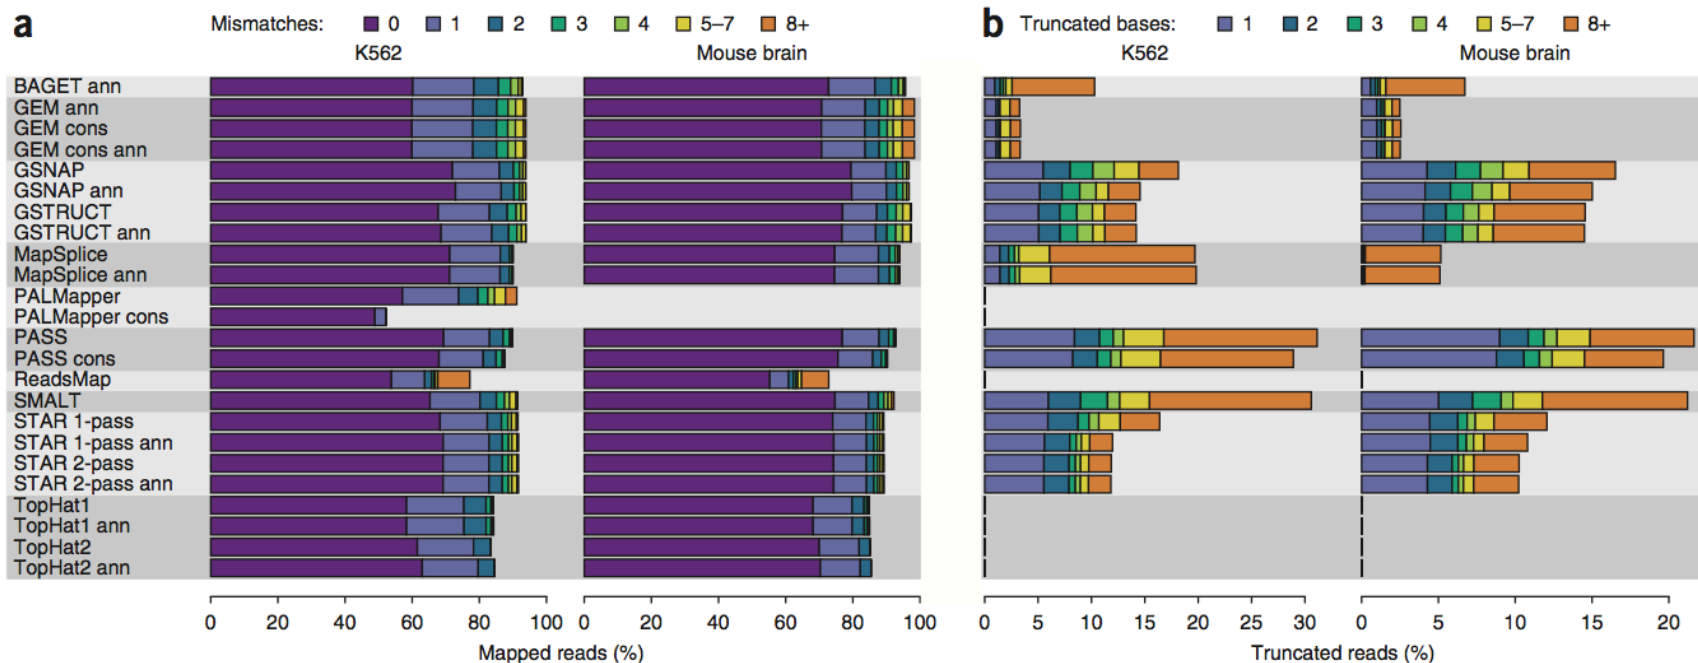
- Designed to align the non-contiguous sequences directly to the reference genome
- Steps:
 - Search Maximal Mappable Prefix (MMP)
 - clustering/stitching/scoring



Comparison of Aligners

- Tradeoff:
 - speed
 - RAM
 - accuracy
 - sensitivity
- <http://www.ecseq.com/support/benchmark>

Performance comparison of RNA-seq mappers

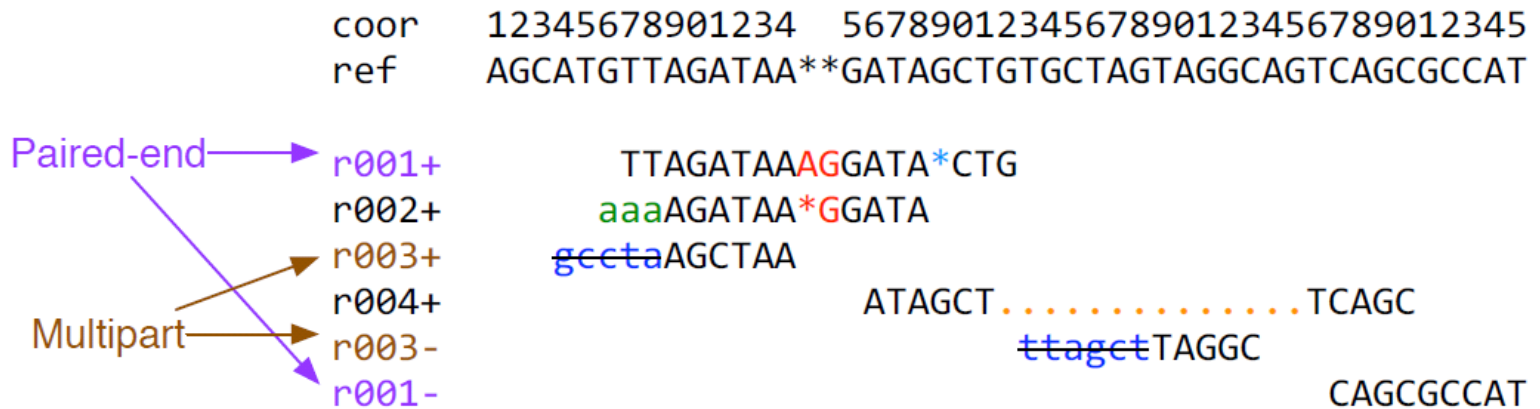




Sequence / Alignment (SAM) files

- **SAM (Sequence Alignment/Map)**
 - Single unified format for storing read alignments to a reference genome
 - Large plain text file
- **BAM (Binary Alignment/Map)**
 - Binary equivalent of SAM
 - Compressed data plus index (bai)
 - Developed for fast processing/indexing


An example of read mapping



Corresponding SAM file

	coord	12345678901234	5678901234567890123456789012345
	ref	AGCATGTTAGATAA**GATAGCTGTGCTAGTAGGCAGTCAGCGCCAT	
Paired-end	r001+	TTAGATAAAGGATA*CTG	
	r002+	aaaAGATAA*GGATA	
	r003+	gcctaAGCTAA	
	r004+	ATAGCT.....TCAGC	
Multipart	r003-	ttagctTAGGC	
	r001-	CAGCGCCAT	
	@SQ	SN:ref LN:45	
Ins & padding	r001	163 ref 7 30 8M2I4M1D3M = 37 39 TTAGATAAAGGATACTA *	
Soft clipping	r002	0 ref 9 30 3S6M1P1I4M * 0 0 AAAAGATAAGGATA *	
	r003	0 ref 9 30 5H6M * 0 0 AGCTAA * NM:i:1	
Splicing	r004	0 ref 16 30 6M14N5M * 0 0 ATAGCTTCAGC *	
Hard clipping	r003	16 ref 29 30 6H5M * 0 0 TAGGC * NM:i:0	
	r001	83 ref 37 30 9M = 7 -39 CAGCGCCAT *	
	QNAME	RNAME	MAPQ?
	FLAG	POS	CIGAR?
		MRNM	ISIZE
		MPOS	SEQ
			QUAL

Slide modified from Heng Li's presentation



CIGAR string - compact representation of an alignment

- M - match or mismatch
- **I** – insertion
- **D** – deletion
- **S** - soft clip
 - Clipped sequences stored in SAM
- **H** - hard clip
 - Clipped sequences not stored in SAM
- N – skipped reference bases, splicing

Match/mismatch, indels

Ref: ACGCA**G**TG--GT

Read: ATGCA-TG**C**AGT

Cigar: 5M**1D**2M**2I**2M

Soft clipping

REF: **A**TCGTGTAACCT**G**ACTAGT**TAA**

READ: **ggg**GTGTAACC-GACTAG**ggg**

Cigar: **3S**8M**1D**5M**4S**

Hard clipping

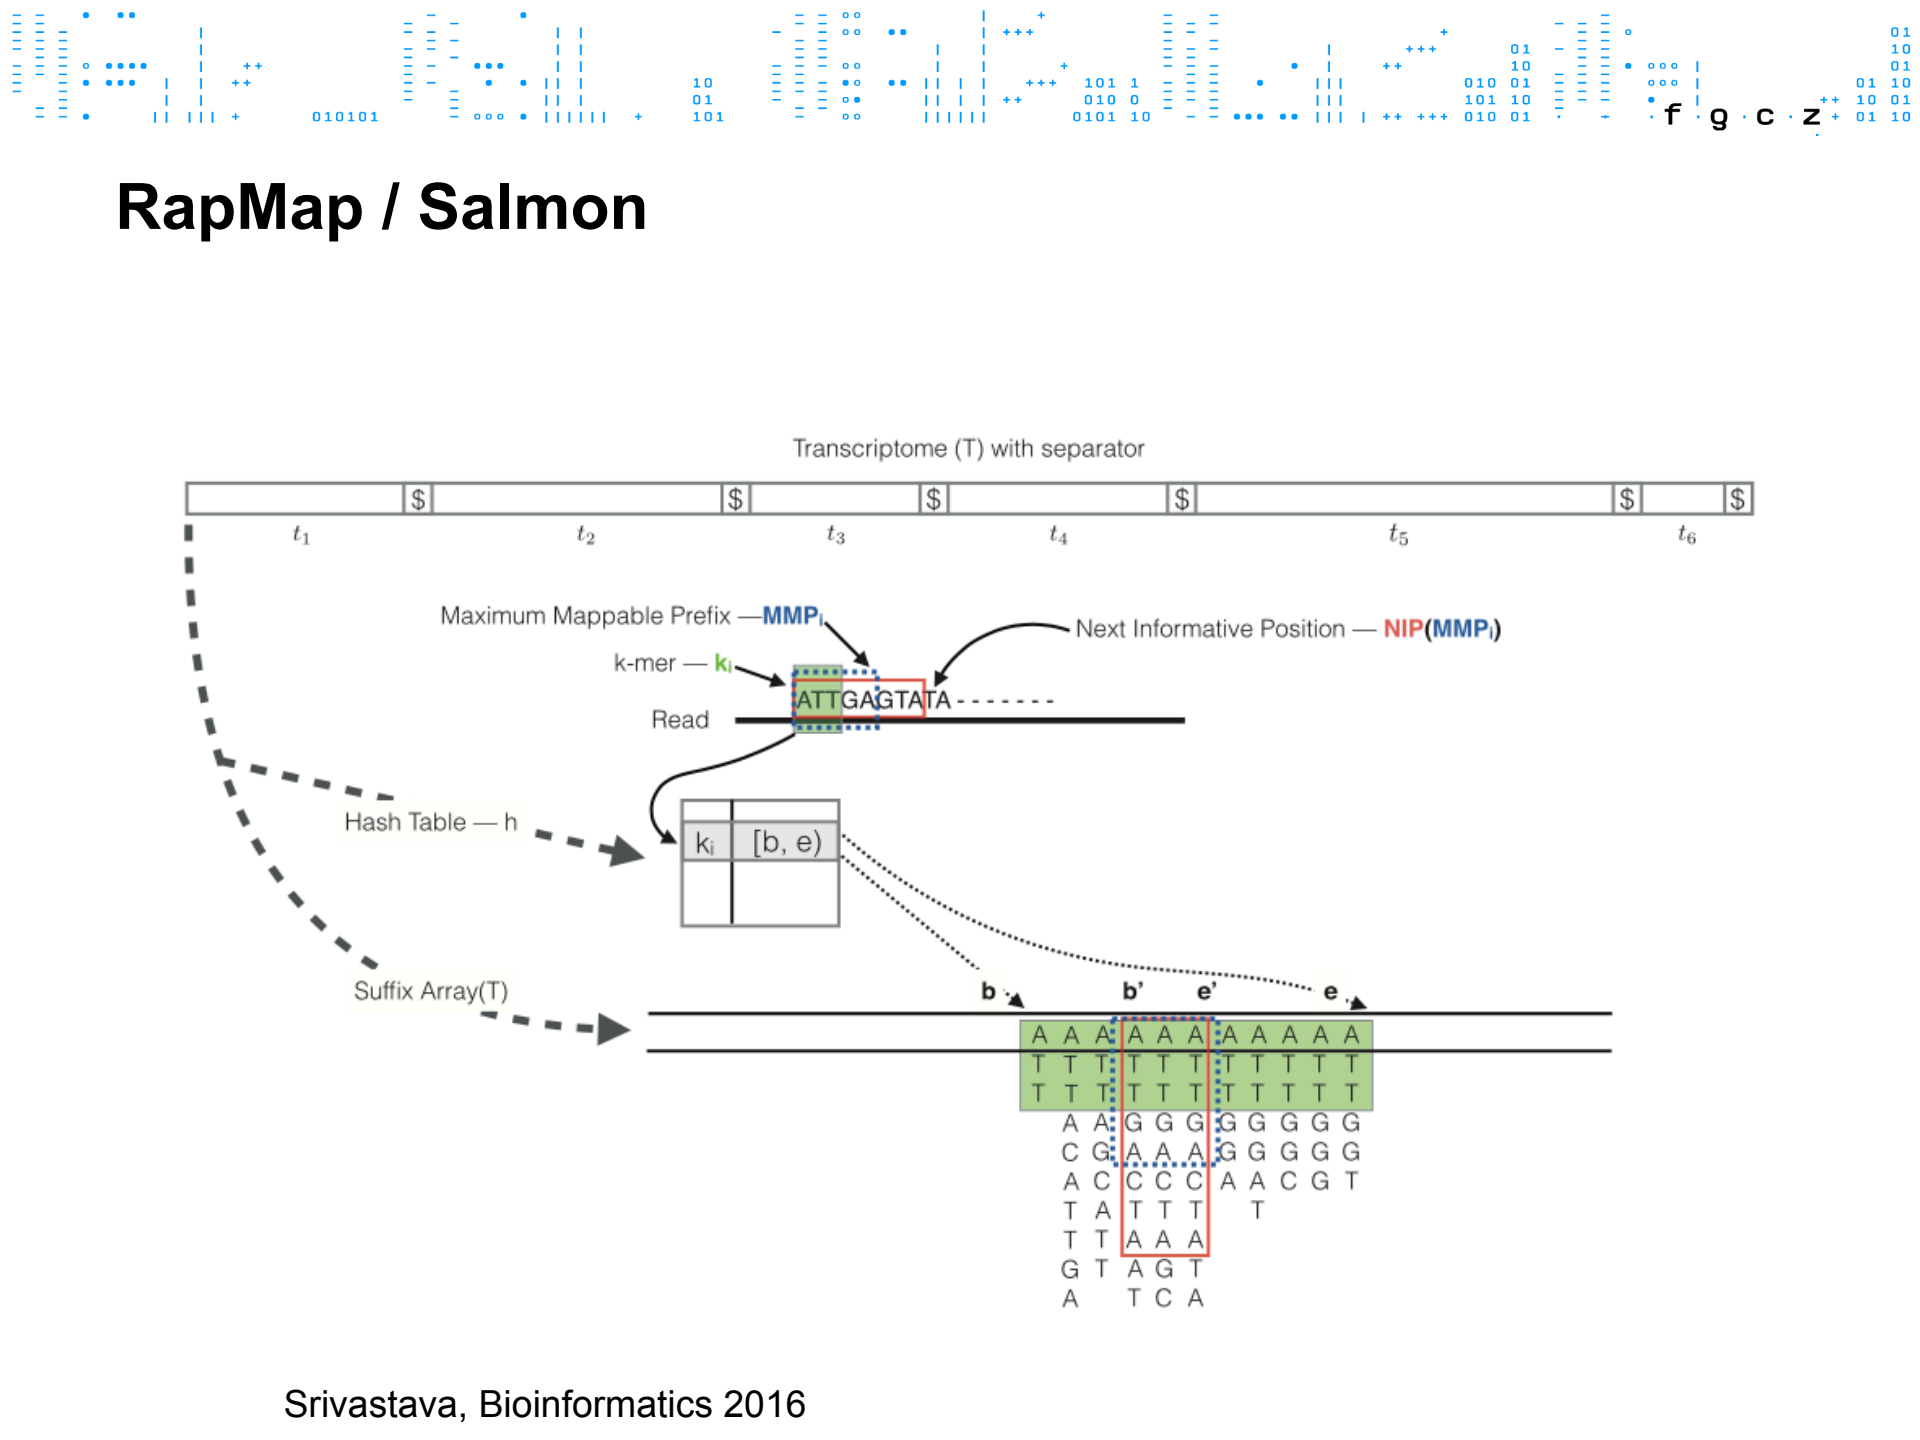
REF: **A**TCGTGTAACCT**G**ACTAGT**TAA**

READ: **ggg**GTGTAACC-GACTAG**ggg**

Cigar: **3H**8M**1D**5M**4H**

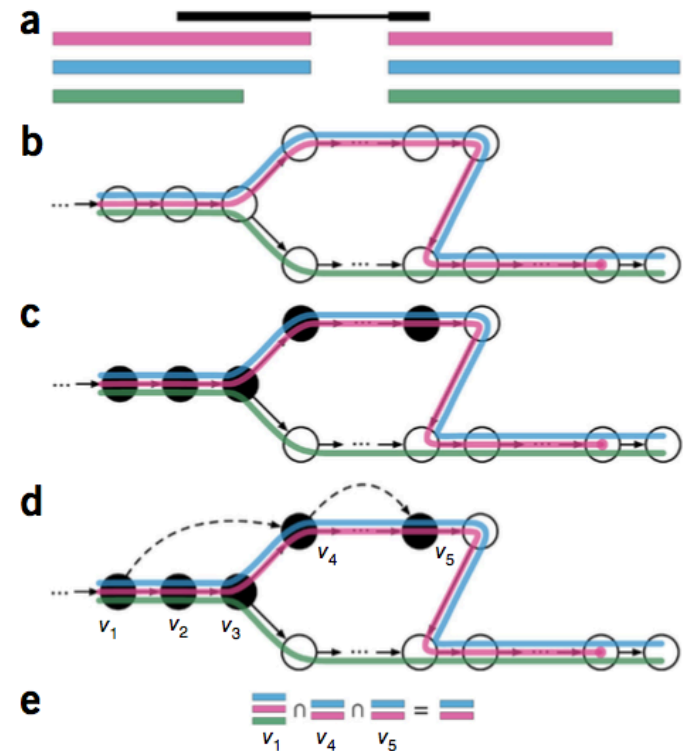
Quasi-Alignments or Pseudo-Alignments


- compare reads only against a hashed transcriptome index
- use only perfect alignments of short read fragments (k-mers, e.g. with k=31).

[illegible]

kallisto pseudo-alignments

- Create every k-mer in the transcriptome (k=31), build de Bruin Graph and color each k-mer
- Preprocess the transcriptome to create the T-DBG






Each k-mer appears in a set of transcripts

The intersection of all sets is our pseudoalignment

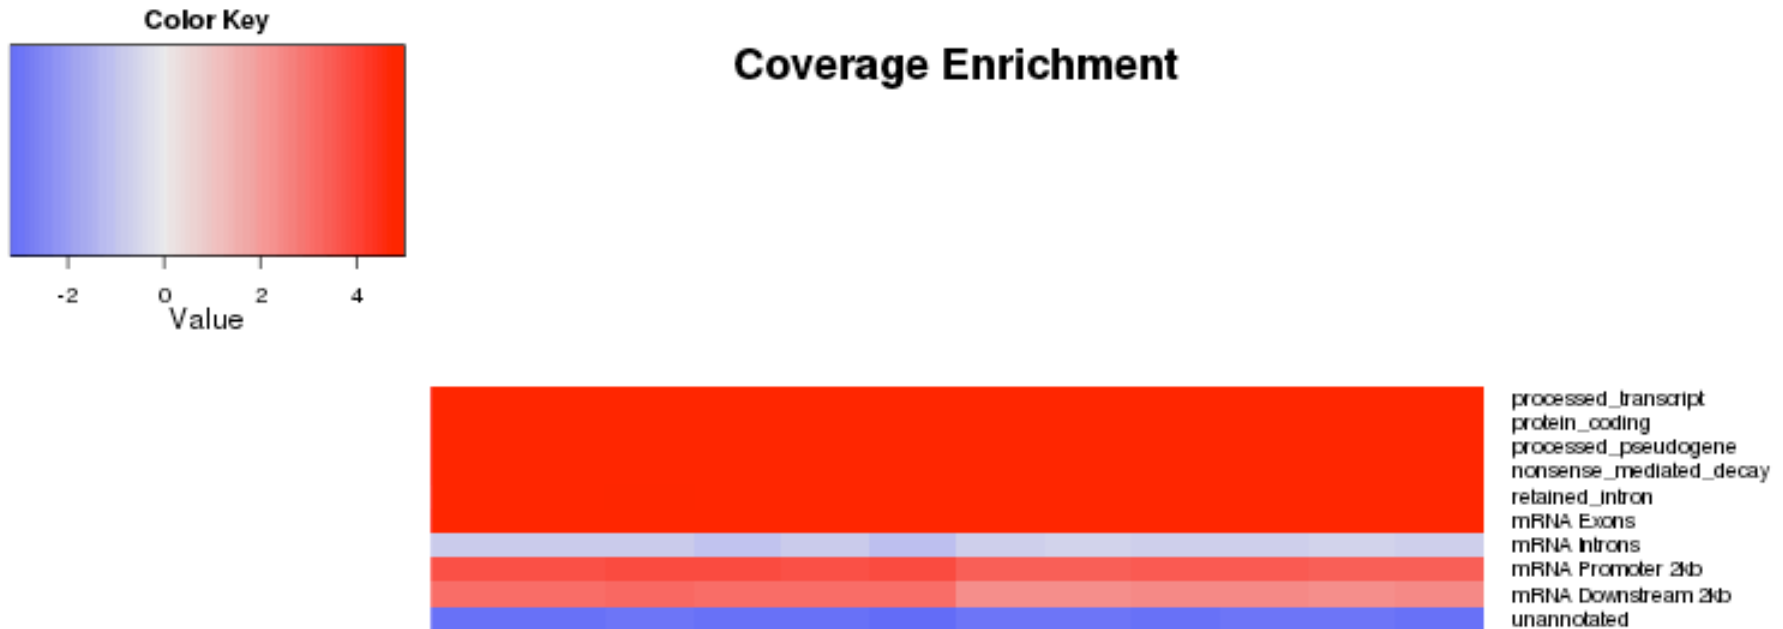
Can jump over k-mers in the T-DBG that provide same information Jumping provides ~8x speedup over checking all k-mers

Mapping QC: Mapping statistics


- How well did my sequence library align to my reference?
- Summary statistics (per read library)
 - % reads with unique alignment
 - % reads with multiple alignment
 - % reads with no alignment
 - % reads properly paired (for paired end libraries)



Mapping QC: Profiling efficiency by transcript annotation



- Transcript annotation: intron, exon, up/down stream, unannotated
- Expression profiling efficiency






Mapping QC: Coverage bias


3' bias

No bias

5' bias



Mapping QC: Junction saturation



- Depth needed for alternative splicing analysis
- All annotated splice junctions are detected - a saturated RNA-seq dataset



RNA-seq mapping QC tools

- RNA-SeQC
 - <https://confluence.broadinstitute.org/display/CGATools/RNA-SeQC>
- EVER-seq (RSeQC)
 - <http://code.google.com/p/rseqc/>
- Qualimap
 - <http://qualimap.bioinfo.cipf.es/>